

Christopher Musco

370 Jay Street, Floor 11 – Brooklyn, NY 11201 – Office 1105

☎ (401) 578 2541 • ✉ cmusco@nyu.edu • 🌐 www.chrismusco.com

Academic Positions

New York University, Tandon School of Engineering <i>Department of Computer Science and Engineering</i> Institute Associate Professor (with tenure) Assistant Professor	New York, NY 2025 - present 2019-2025
Brown University <i>Visiting Scholar, Institute for Computational and Experimental Research in Mathematics</i>	Providence, RI Spring 2026
University of California, Berkeley <i>Visiting Researcher, Simons Institute for the Theory of Computing</i>	Berkeley, CA Fall 2025
Microsoft Research <i>Visiting Researcher, Machine Learning and Optimization Group</i>	Redmond, WA Summer 2020
Princeton University <i>Research Instructor, Department of Computer Science</i>	Princeton, NJ 2018 – 2019
Massachusetts Institute of Technology <i>Research and Teaching Assistant, Theory of Computation Group</i>	Cambridge, MA 2013 – 2018
Yahoo Labs <i>Research Intern, Scalable Machine Learning Group</i>	New York, NY Summer 2015

Research Areas

theory of algorithms • randomized methods • numerical linear algebra • sketching and streaming • high-dimensional data • machine learning and ML for science • active learning and function approximation

Education

Massachusetts Institute of Technology <i>Ph.D. in Computer Science</i> Advisor: Jonathan Kelner Ph.D. Thesis: Faster Linear Algebra for Data Analysis and Machine Learning M.S. Thesis: Dimensionality Reduction for Sparse and Structured Matrices	Cambridge, MA 2013 – 2018
Yale University <i>B.S. in Applied Mathematics, B.S. in Computer Science</i>	New Haven, CT 2008 – 2012

Publications

Authors appear in alphabetical order, in the tradition of mathematics and theoretical computer science. Any exceptions are marked with a *, which indicates that authors are ordered by contribution.

Almost Navigable Graphs. Pratyush Avi, Christopher Musco. *VecDB Workshop at VLDB 2026.*

Quantizing Sampling-Based Inner Product Sketches*. Gabriel M. Silva, Aécio Santos, Majid Daliri, Juliana Freire, Christopher Musco. *VecDB Workshop at VLDB* 2026.

Nearly Instance Optimal Sparse Matrix Approximation from Matrix-Vector Products. Christopher Musco, Indu Ramesh. *European Symposium on Algorithms (ESA)* 2026.

Spectral density estimation for normal matrices. Cameron Musco, Christopher Musco, Rikhav Shah, John Urschel, Nicholas West. Preprint 2026.

Query Efficient Structured Matrix Learning. Noah Amsel, Pratyush Avi, Tyler Chen, Feyza Duman Keles, Chinmay Hegde, Cameron Musco, Christopher Musco, David Persson. *Conference on Learning Theory (COLT)* 2026. Preliminary version in Theory of AI for Scientific Computing Workshop at COLT 2025.

Randomized Randomized block-Krylov subspace methods for low-rank approximation of matrix functions*. David Persson, Tyler Chen, Christopher Musco. *Linear Algebra and Its Applications (LAA)* 2026.

The Polar Express: Optimal Matrix Sign Methods and Their Application to the Muon Algorithm*. Noah Amsel, David Persson, Christopher Musco, Robert Gower. *International Conference on Learning Representations (ICLR)* 2026. Outstanding Paper Award Honorable Mention (top 3 of 19,525 submissions).

PolySHAP: Extending KernelSHAP with Interaction-Informed Polynomial Regression*. Fabian Fumagalli, R. Teal Witter, Christopher Musco. *International Conference on Learning Representations (ICLR)* 2026.

An Exact Algorithm for the Unanimous Vote Problem. Feyza Duman Keles, Lisa Hellerstein, Kunal Marwaha, Christopher Musco, Xincheng Yang. *SIAM Symposium on Simplicity in Algorithms (SOSA)* 2026.

Does block size matter in randomized block Krylov low-rank approximation?. Tyler Chen, Ethan N. Epperly, Raphael A. Meyer, Christopher Musco, Akash Rao. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2026.

Efficiently Constructing Sparse Navigable Graphs. Alex Conway, Laxman Dhulipala, Martin Farach-Colton, Rob Johnson, Ben Landrum, Christopher Musco, Yarin Shechter, Torsten Suel, Richard Wen. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2026.

Quasi-Optimal Hierarchically Semi-Separable Matrix Approximation. Noah Amsel, Tyler Chen, Feyza Duman Keles, Diana Halikias, Cameron Musco, Christopher Musco, David Persson. *SIAM Journal on Matrix Analysis and Applications (SIMAX)* 2026.

Fixed-Sparsity Matrix Approximation from Matrix-Vector Products. Noah Amsel, Tyler Chen, Diana Halikias, Feyza Duman Keles, Cameron Musco, Christopher Musco. *SIAM Journal on Matrix Analysis and Applications (SIMAX)* 2026.

Distance Adaptive Beam Search for Provably Accurate Graph-Based Nearest Neighbor Search. Yousef Al-Jazzazi, Haya Diwan, Jinrui Gou, Cameron Musco, Christopher Musco, Torsten Suel. *Conference on Neural Information Processing Systems (NeurIPS)* 2025.

Regression-adjusted Monte Carlo Estimators for Shapley Values and Probabilistic Values*. R. Teal Witter, Yurong Liu, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2025.

Sharper Bounds for Chebyshev Moment Matching with Applications to Differential Privacy and Beyond. Cameron Musco, Christopher Musco, Lucas Rosenblatt, Apoorv Vikram Singh. *Conference on Learning Theory (COLT)* 2025.

Coupling without Communication and Drafter-Invariant Speculative Decoding. Majid Daliri, Christopher Musco, Ananda Theertha Suresh. *International Symposium on Information Theory (ISIT)* 2025.

Kernel Banzhaf: A Fast and Robust Estimator for Banzhaf Values*. Yurong Liu, R. Teal Witter, Flip Korn, Tarfah Alrashed, Dimitris Paparas, Christopher Musco, Juliana Freire. Preprint 2025.

Matrix Product Sketching via Coordinated Sampling. Majid Daliri, Juliana Freire, Danrong Li, Christopher Musco. *International Conference on Learning Representations (ICLR)* 2025.

Provably Accurate Shapley Value Estimation via Leverage Score Sampling. Christopher Musco, R. Teal Witter. *International Conference on Learning Representations (ICLR)* 2025. Selected as spotlight paper (5.1% of submissions).

Improved Spectral Density Estimation via Explicit and Implicit Deflation. Rajarshi Bhattacharjee, Rajesh Jayaram, Cameron Musco, Christopher Musco, Archan Ray. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2025.

Near-optimal Hierarchical Matrix Approximation from Matrix-vector Products. Tyler Chen, Diana Halikias, Feyza Duman Keles, Cameron Musco, Christopher Musco, David Persson. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2025.

Faster Linear Systems and Matrix Norm Approximation via Multi-level Sketched Preconditioning. Michał Dereziński, Christopher Musco, Jiaming Yang. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2025.

Algorithm-agnostic low-rank approximation of operator monotone matrix functions*. David Persson, Raphael A. Meyer, Christopher Musco. *SIAM Journal on Matrix Analysis and Applications (SIMAX)* 2025.

Navigable Graphs for High-Dimensional Nearest Neighbor Search: Constructions and Limits. Haya Diwan, Jinrui Gou, Cameron Musco, Christopher Musco, Torsten Suel. *Conference on Neural Information Processing Systems (NeurIPS)* 2024.

Near-Optimality Guarantees for Approximating Rational Matrix Functions by the Lanczos Method. Noah Amsel, Tyler Chen, Anne Greenbaum, Cameron Musco, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2024. Selected as spotlight paper (2.1% of submissions).

Benchmarking Estimators for Natural Experiments: A Novel Dataset and a Doubly Robust Algorithm*. R. Teal Witter, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2024.

Faster Spectral Density Estimation and Sparsification in the Nuclear Norm. Yujia Jin, Ishani Karmarkar, Christopher Musco, Aaron Sidford, Apoorv Vikram Singh. *Conference on Learning Theory (COLT)* 2024.

Agnostic Active Learning of Single Index Models with Linear Sample Complexity*. Aarshvi Gajjar, Wai Ming Tai, Xingyu Xu, Chinmay Hegde, Christopher Musco, Yi Li. *Conference on Learning Theory (COLT)* 2024.

Sampling Methods for Inner Product Sketching . Majid Daliri, Juliana Freire, Christopher Musco, Aécio Santos, Haoxiang Zhang. *Proceedings of the VLDB Endowment (VLDB)*. 2024.

Improved Active Learning via Dependent Leverage Score Sampling*. Atsushi Shimizu, Xiaouu Cheng, Christopher Musco, Jonathan Weare. *International Conference on Learning Representations (ICLR)* 2024. Invited for full oral presentation (1.2% of submissions).

Improved Bounds for Agnostic Active Learning of Single Index Models. Aarshvi Gajjar, Xingyu Xu, Christopher Musco, Chinmay Hegde. *NeurIPS Workshop on Adaptive Experimental Design and Active Learning in the Real World* 2023.

A Simple and Practical Method for Reducing the Disparate Impact of Differential Privacy*. Lucas Rosenblatt, Julia Stoyanovich, Christopher Musco. *AAAI Conference on Artificial Intelligence (AAAI)* 2024.

On the Unreasonable Effectiveness of Single Vector Krylov Methods for Low-Rank Approximation. Raphael A. Meyer, Cameron Musco, Christopher Musco. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2024.

Simple Analysis of Priority Sampling. Majid Daliri, Juliana Freire, Christopher Musco, Aécio Santos, Haoxiang Zhang. *SIAM Symposium on Simplicity in Algorithms (SOSA)* 2024.

Structured Semidefinite Programming for Recovering Structured Preconditioners. Arun Jambulapati, Jerry Li, Christopher Musco, Aaron Sidford, Kevin Tian. *Conference on Neural Information Processing Systems*

(NeurIPS) 2023.

Efficient Block Approximate Matrix Multiplication*. Chuhan Yang, Christopher Musco. *European Symposium on Algorithms (ESA)* 2023.

Moments, Random Walks, and Limits for Spectrum Approximation. Yujia Jin, Christopher Musco, Apoorv Singh, Aaron Sidford. *Conference on Learning Theory (COLT)* 2023.

Dimensionality Reduction for General KDE Mode Finding. Xinyu Luo, Christopher Musco, Cas Widdershoven. *International Conference on Machine Learning (ICML)* 2023.

Weighted Minwise Hashing Beats Linear Sketching for Inner Product Estimation. Aline Bessa, Majid Daliri, Juliana Freire, Cameron Musco, Christopher Musco, Aécio Santos, Haoxiang Zhang. *ACM Symposium on Principles of Database Systems (PODS)* 2023.

Low-memory Krylov Subspace Methods for Optimal Rational Matrix Function Approximation. Tyler Chen, Anne Greenbaum, Cameron Musco, Christopher Musco. *SIAM Journal on Matrix Analysis and Applications (SIMAX)* 2023.

Active Learning for Single Neuron Models with Lipschitz Non-Linearities. Aarshvi Gajjar, Chinmay Hegde, Christopher Musco. *International Conference on Artificial Intelligence and Statistics (AISTATS)* 2023. Preliminary version at NeurIPS 2022 Workshop on The Symbiosis of Deep Learning and Differential Equations II (DLDE 2022).

Near-Linear Sample Complexity for L_p Polynomial Regression. Raphael A. Meyer, Cameron Musco, Christopher Musco, David P. Woodruff, Samson Zhou. *ACM-SIAM Symposium on Discrete Algorithms (SODA)*. 2023.

A Tight Analysis of Hutchinson's Diagonal Estimator. Prathamesh Dharangutte, Christopher Musco. *SIAM Symposium on Simplicity in Algorithms (SOSA)* 2023.

Active Linear Regression for ℓ_p Norms and Beyond. Cameron Musco, Christopher Musco, David P. Woodruff, Taisuke Yasuda. *IEEE Symposium on Foundations of Computer Science (FOCS)* 2022.

How to Quantify Polarization in Models of Opinion Dynamics. Christopher Musco, Indu Ramesh, Johan Ugander, R. Teal Witter. *International Workshop on Mining and Learning with Graphs (at KDD)* 2022.

Chemically-informed data-driven optimization (ChIDDO): Leveraging physical models and Bayesian learning to accelerate chemical research*. Daniel Frey, Juhee Shin, Christopher Musco, Miguel A. Modestino. *Reaction Chemistry & Engineering* 2022.

Streaming Approach to In Situ Selection of Key Time Steps for Time-Varying Volume Data*. Mengxi Wu, Yi-Jen Chiang, Christopher Musco. *EuroVis* 2022.

Sublinear Time Spectral Density Estimation. Vladimir Braverman, Aditya Krishnan, Christopher Musco. *ACM Symposium on Theory of Computing (STOC)* 2022.

A Sketch-based Index for Correlated Dataset Search*. Aécio Santos, Aline Bessa, Christopher Musco, Juliana Freire. *IEEE 38th International Conference on Data Engineering (ICDE)* 2022.

Fast Regression for Structured Inputs. Raphael A. Meyer, Cameron Musco, Christopher Musco, David P. Woodruff, Samson Zhou. *International Conference on Learning Representations (ICLR)* 2022.

Active Linear Regression in the Online Setting via Leverage Score Sampling. Harvineet Singh, Christopher Musco, Rumi Chunara. *ICML Workshop on Adaptive Experimental Design and Active Learning in the Real World* 2022.

Error Bounds for Lanczos-based Matrix Function Approximation. Tyler Chen, Anne Greenbaum, Cameron Musco, Christopher Musco. *SIAM Journal on Matrix Analysis and Applications (SIMAX)* 2022.

Dynamic Trace Estimation. Prathamesh Dharangutte, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2021.

Finding an Approximate Mode of a Kernel Density Estimate. Jasper C.H. Lee, Jerry Li, Christopher Musco, Jeff M. Phillips, Wai Ming Tai. *European Symposium on Algorithms (ESA)* 2021.

Correlation Sketches for Approximate Join-Correlation Queries*. Aécio Santos, Aline Bessa, Fernando Chirigati, Christopher Musco, Juliana Freire. *ACM Special Interest Group on Management of Data Conference (SIGMOD)* 2021.

Public Transport Planning: When Transit Network Connectivity Meets Commuting Demand*. Sheng Wang, Yuan Sun, Christopher Musco, Zhifeng Bao. *ACM Special Interest Group on Management of Data Conference (SIGMOD)* 2021.

Simple Heuristics Yield Provable Algorithms for Masked Low Rank Approximation. Cameron Musco, Christopher Musco, David Woodruff. *Innovations in Theoretical Computer Science (ITCS)* 2021.

Hutch++: Optimal Stochastic Trace Estimation. Raphael A. Meyer, Cameron Musco, Christopher Musco, David P. Woodruff. *SIAM Symposium on Simplicity in Algorithms (SOSA)* 2021.

Graph Learning for Inverse Landscape Genetics. Prathamesh Dharangutte, Christopher Musco. *AAAI Conference on Artificial Intelligence (AAAI)* 2021. Short version in NeurIPS 2020 AI for Earth Sciences workshop.

Fourier Sparse Leverage Scores and Approximate Kernel Learning. Tamás Erdélyi, Cameron Musco, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2020. Invited for spotlight presentation (4% of submissions).

Near Optimal Linear Algebra in the Online and Sliding Window Models. Vladimir Braverman, Petros Drineas, Cameron Musco, Christopher Musco, Jalaj Upadhyay, David P. Woodruff, Samson Zhou. *IEEE Symposium on Foundations of Computer Science (FOCS)* 2020.

The Statistical Cost of Robust Kernel Hyperparameter Tuning. Raphael A. Meyer, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2020.

Low-Rank Toeplitz Matrix Estimation via Random Ultra-Sparse Rulers. Hannah Lawrence, Jerry Li, Cameron Musco, Christopher Musco. *International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*. 2020.

Sample Efficient Toeplitz Covariance Estimation. Yonina Eldar, Jerry Li, Cameron Musco, Christopher Musco. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2020.

Fast and Space Efficient Spectral Sparsification in Dynamic Streams. Michael Kapralov, Aida Mousavifar, Cameron Musco, Christopher Musco, Navid Nouri, Aaron Sidford. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2020.

Analyzing the Impact of Filter Bubbles on Social Network Polarization. Uthsav Chitra, Christopher Musco. *ACM International Conference on Web Search and Data Mining (WSDM)* 2020. Preliminary version in KDD Workshop on Issues of Sentiment Discovery and Opinion Mining.

Projection-Cost-Preserving Sketches: Proof Strategies and Constructions. Cameron Musco, Christopher Musco. *Technical Report* 2020.

A Universal Sampling Method for Reconstructing Signals with Simple Fourier Transforms. Haim Avron, Michael Kapralov, Cameron Musco, Christopher Musco, Ameya Velingker, Amir Zandieh. *ACM Symposium on Theory of Computing (STOC)* 2019.

Inferring Networks from Random Walk-based Node Similarities. Jeremy Hoskins, Cameron Musco, Christopher Musco, Charalampos Tsourakakis. *Conference on Neural Information Processing Systems (NeurIPS)* 2018.

Eigenvector Computation and Community Detection in Asynchronous Gossip Models. Frederik Mallmann-Trenn, Cameron Musco, Christopher Musco. *International Colloquium on Automata, Languages, and Programming (ICALP)* 2018.

Minimizing Polarization and Disagreement in Social Networks. Cameron Musco, Christopher Musco, Charalampos Tsourakakis. *The Web Conference (WWW)* 2018.

Stability of the Lanczos Method for Matrix Function Approximation. Cameron Musco, Christopher Musco, Aaron Sidford. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2018.

Recursive Sampling for the Nyström Method. Cameron Musco, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2017.

Random Fourier Features for Kernel Ridge Regression: Approximation Bounds and Statistical Guarantees. Haim Avron, Michael Kapralov, Cameron Musco, Christopher Musco, Ameya Velingker, Amir Zandieh. *International Conference on Machine Learning (ICML)* 2017.

Input Sparsity Time Low-Rank Approximation via Ridge Leverage Score Sampling. Michael B. Cohen, Cameron Musco, Christopher Musco. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2017.

Determining Tournament Payout Structures for Daily Fantasy Sports. Christopher Musco, Maxim Sviridenko, Justin Thaler. *SIAM Algorithm Engineering and Experiments (ALENEX)* 2017.

Principal Component Projection Without Principal Component Analysis. Roy Frostig, Cameron Musco, Christopher Musco, Aaron Sidford. *International Conference on Machine Learning (ICML)* 2016.

Randomized Block Krylov Methods for Stronger and Faster Approximate Singular Value Decomposition. Cameron Musco, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2015. Invited for full oral presentation (1 of 15, < 1% of submissions). Also presented at the 2016 Copper Mountain Conference on Iterative Methods.

Dimensionality Reduction for K-Means Clustering and Low Rank Approximation. Michael B. Cohen, Samuel Elder, Cameron Musco, Christopher Musco, Madalina Persu. *ACM Symposium on Theory of Computing (STOC)* 2015.

Principled Sampling for Anomaly Detection. Brendan Juba, Fan Long, Christopher Musco, Stelios Sidiropoulos-Douskos, Martin Rinard. *Network and Distributed System Security Symposium (NDSS)*. 2015.

Uniform Sampling for Matrix Approximation. Michael B. Cohen, Yin Tat Lee, Cameron Musco, Christopher Musco, Richard Peng, Aaron Sidford. *Innovations in Theoretical Computer Science (ITCS)* 2015.

Single Pass Spectral Sparsification in Dynamic Streams. Michael Kapralov, Yin Tat Lee, Cameron Musco, Christopher Musco, Aaron Sidford. *IEEE Symposium on Foundations of Computer Science (FOCS)* 2014, Special Issue SIAM Journal on Computing 2017.

Funding

National Science Foundation, Future CoRe

Graph-based Vector Databases, from Theory to Practice. 2026 - 2030
 PI (w/ Alex Conway, Cornell, Martin Farach-Colton, NYU, Laxman Dhulipala, UMD, Michael Bender, Stony Brook), \$1,000,000

Google

Scalable Chebyshev Moment Estimation for Differentially Private Synthetic Data 2026
 PI (w/ Lucas Rosenblatt, NYU), \$20,000 in Google Cloud Credit

National Science Foundation, AF Small

The Query Complexity of Linear Algebra. 2024 - 2027
 PI (w/ Cameron Musco, UMass), \$600,000

National Science Foundation, CAREER Award

Frontiers in Matrix Sketching. 2021 - 2027
 PI, \$562,224

National Science Foundation, III Medium

Dataset Search and Ranking for Data Augmentation and Explanation. 2021 - 2025
 Co-PI (w/ Juliana Freire, NYU), \$1,093,195

DOE Office of Science Grant*Leverage Score Sampling for Parametric PDEs.*

2021 - 2023

PI (w/ Chinmay Hegde, NYU), \$300,000

National Science Foundation/Computing Research Association*CIFellow Fellow Award for Dr. Aline Bessa.*

2021-2022

PI, \$263,250

Current Graduate Student and Postdoc Advising

Pratyush Avi. Ph.D., New York University.	2024 - present
Noah Amsel. Ph.D., New York University (with Prof. Joan Bruna).	2022 - present
Majid Daliri. Ph.D., New York University.	2022 - present
Venetia Pliatsika. Ph.D., New York University.	2026 - present
Apoorv Singh. Ph.D., New York University.	2020 - present

Past Advising

Ph.D. Students and Postdocs

Feyza Duman Keles. Ph.D., New York University.	2023 - 2026
Aarsvhi Gajjar. Ph.D., New York University.	2021 - 2025
Teal Witter. Ph.D., New York University.	2020 - 2025
Tyler Chen. Postdoc, New York University.	2022 - 2024
Raphael A. Meyer. Ph.D., New York University.	2019 - 2024
Chuhan Yang. Ph.D., New York University.	2020 - 2024
David Persson. Visiting Ph.D. student.	2023
Aline Bessa. Postdoc, New York University.	2021

Master's Students

Rishi Kaushik Rao, M.S., New York University.	2025 - 2026
Shubham Rastogi, M.S., New York University.	2025
Akash Rao, M.S., New York University.	2023 - 2024
Siddharth Sagar. M.S., New York University.	2023
Yajur Ahuja. M.S., New York University.	2023
Sundeeep Chenreddy. M.S., New York University.	2023
Atsushu Shimuzu. M.S., New York University.	2022 - 2023
Danrong Li. M.S., New York University.	2022 - 2023
Indu Ramesh. M.S., New York University.	2020 - 2022
Mengxi Wu. M.S., New York University.	2020 - 2022
Xinyu Luo. M.S., New York University.	2020 - 2022
Prathamesh Dharangutte. M.S., New York University.	2019 - 2021
John Shin. M.S., New York University.	2019 - 2020

I have advised over a dozen undergraduate guided studies and research projects, as well as several research projects involving high school students.

Ph.D. Thesis Committees

Sivan Ding, NYU	
Kshiteej Sheth, EPFL	<i>Defended 2026</i>
Xiaoou Cheng, NYU	<i>Defended 2026</i>
Lucas Rosenblatt, NYU	<i>Defended 2026</i>
Oğuzhan Kubilay Ülger, NYU	<i>Defended 2025</i>
Hanna Komlós, NYU	<i>Defended 2025</i>
Bruno Coelho, NYU	<i>Defended 2025</i>
Haoxiang Zhang, NYU	<i>Defended 2024</i>
Aécio Santos, NYU	<i>Defended 2024</i>
Ishan Agarwal, NYU	<i>Defended 2023</i>
Yiming Zhang , NYU	<i>Defended 2022</i>
Francis Williams, NYU	<i>Defended 2021</i>

Teaching

Algorithmic Machine Learning and Data Science (NYU CS-GY 6763)	
<i>Course Instructor</i>	<i>Fall 2019-2023, Spring 2025</i>
Recent Developments in Algorithm Design (NYU CSCI-GA 3033/CS-GY 9223)	
<i>Co-Instructor</i>	<i>Spring 2025</i>
Machine Learning, Graduate (NYU CS-GY 6923)	
<i>Course Instructor</i>	<i>Spring 2022, 2023, Fall 2024</i>
Introduction to Machine Learning (NYU CS-UY 4563)	
<i>Course Instructor</i>	<i>Spring 2020</i>
Advanced Algorithm Design (Princeton COS 521)	
<i>Course Instructor</i>	<i>Fall 2018</i>
Advanced Algorithms (MIT 6.854/18.415)	
<i>Teaching Assistant</i>	<i>Spring 2016</i>
Technical Communication Skills for Graduate Students (MIT 6.S977)	
<i>Workshop Leader</i>	<i>Spring 2016</i>

Mentorship and Outreach

SIAM/MathWorks Math Modeling Challenge	
<i>Problem Development Committee, Director of Technical Computing Judging</i>	<i>2018 - present</i>
Support SIAM's annual math modeling competition, which gives high school students an opportunity to apply mathematics, statistics, and data science to a real-world problem. Serve on committee in charge of writing practice and contest problems. Lead judging and award presentation for student Technical Computing prizes.	
MIT Graduate Communication Lab	
<i>Communication Adviser</i>	<i>2015 - 2018</i>
Helped to establish an EECS Communication Lab to provide over 700 graduate students with support in academic writing, oral presentation, job applications, and generally, in communicating effectively. Offered feedback, guidance, and coaching to individual students during weekly office hours.	

Invited Talks

Structured Matrix Approximation from Matrix-Vector Products

ICERM at Brown University

Feb. 2026

Rensselaer Polytechnic Institute

Oct. 2024

Randomized Numerical Linear Algebra: Highlights and Future Directions

ICERM at Brown University

Jan. 2026

Simons Institute at Berkeley

Nov. 2025

Navigability and Graph-based Vector Search

Princeton University, Theory Lunch

Feb. 2026

CUNY Graduate Center, Holiday Lecture

Dec. 2025

Northwestern University, IDEAL Workshop

Nov. 2025

Solving Linear Systems: Recent Developements and Open Questions

Simons Institute at Berkeley

Sept. 2025

Beam Search, Navigable Graphs, and Nearest Neighbor Search

Microsoft Research

July 2025

Mathematically Modeling Anything and Everything in Python

Mu Alpha Theta National Convention

July 2025

On the Effectiveness of Single Vector Krylov Methods for Low-rank Approximation

SIAM Conference on Applied Linear Algebra

May 2024

Foundations of Computational Mathematics (FoCM)

June 2023

The Lanczos Method, Matrix Functions, and the Quest for Optimality

Johns Hopkins, Applied Math Seminar

Feb. 2024

Tutorial on Randomized Numerical Linear Algebra (invited panelist)

Neural Information Processing Systems (NeurIPS)

Dec. 2023

Robust Active Learning via Leverage Score Sampling

One World Mathematics of Information, Data, and Signals (1W-MINDS) Seminar

Oct. 2023

Canadian Mathematics Society, Winter Meeting

Dec. 2022

Tutorial on Matrix Sketching

BIRS Workshop on Perspectives on Matrix Computations

Mar. 2023

Integrating Education and Research in Early Career Funding Proposals

Joint Mathematics Meetings, SIAM Special Session on Education

Jan. 2023

Iterative Sampling Methods for Matrix Approximation

University of Michigan, Randomized Numerical Linear Algebra Seminar

Oct. 2022

Noisy Implicit Methods for Matrix Functions

SIAM Annual Meeting

July 2022

Sublinear Time Spectral Density Estimation

Queens College CUNY Computer Science Colloquium

Mar. 2023

National University of Singapore, Workshop on Algorithms and Foundations for Data Science

June 2022

Oxford University, Computational Mathematics and Applications Seminar

Feb. 2022

University of British Columbia, Applied Math Seminar

Jan. 2022

Workshop on Algorithms for Large Data

Aug. 2021

Cornell University, Scientific Computing and Numerics Seminar	Mar. 2021
Optimal Stochastic Trace Estimation	
New York University, Numerical Analysis and Scientific Computing Seminar	Apr. 2022
University of Michigan, Randomized Numerical Linear Algebra Seminar	Nov. 2021
Stanford University, Operations Research Seminar	Oct. 2020
A TCS Perspective on Iterative Eigensolvers	
Complexity of Matrix Computations Seminar	Oct. 2021
Mathematical Modeling and Coding in the Classroom	
Mu Alpha Theta National Convention	July 2021
Randomized Algorithms for Probing the Spectrum of a Matrix	
George Mason University, AI Tea Series	Apr. 2021
Importance Sampling for Approximating High-Dimensional Kernel Matrices	
Emory University, Numerical Analysis and Scientific Computing Seminar	Feb. 2021
Structured Covariance Estimation	
Oxford Statistics and Machine Learning in Finance Seminar	June 2022
Yahoo! Research, Seminar	Apr. 2020
Johns Hopkins University, Theory Seminar	Oct. 2019
From Randomized Linear Algebra to Randomized Functional Analysis	
SIAM Mathematics of Data Science Conference	May 2020
University of Utah, Data Science Seminar	Jan. 2020
Rutgers University, DIMACS RandNLA Workshop	Sept. 2019
Learning Signals with Simple Fourier Transforms	
Columbia University, Theory Seminar	Oct. 2018
UC Berkeley, Simons Institute Workshop	Sept. 2018
Princeton University, Theory Seminar	Sept. 2018
The Lanczos Method in Data Science: Challenges and Opportunities	
SIAM Conference on Parallel Processing for Scientific Computing, Minisymposium	Feb. 2020
Georgetown University, CS Seminar	Feb. 2018
Unifying Random Features and Leverage Scores for Kernel Matrix Approximation	
Institute for Advanced Study, Computer Science and Discrete Mathematics Seminar	Feb. 2018
Harvard University, guest lecture in "Sketching Algorithms for Big Data"	Nov. 2017
Microsoft Research, Redmond, Research Seminar	Sept. 2017
Sketching and Sampling Methods for Large Scale Linear Algebra	
Twitter, Cambridge, Research Seminar	July 2017
Recursive Sampling for the Nyström Method	
New England Machine Learning Day	May 2017
Introduction to Linear Sketching	
MIT, Theory Retreat	Sept. 2016
Ridge Leverage Score Sampling	
National Institute of Informatics, Shonan Meeting	July 2016
MIT, Algorithms & Complexity Seminar	July 2016

University of Utah, Data Group Meeting	<i>Jan. 2016</i>
Large Scale Column Subset Selection	
Yahoo Labs, Science Week	<i>July 2015</i>
Dimensionality Reduction for k-Means Clustering	
IBM T.J. Watson, Mathematical Sciences Research Seminar	<i>Aug. 2015</i>
MIT, Theory Lunch	<i>June 2014</i>
Uniform Sampling for Matrix Approximation	
Yahoo Labs, Scalable Machine Learning Research Seminar	<i>June 2015</i>
Single Pass Spectral Sparsification in Dynamic Streams	
MIT, Annual Sublinear Algorithms Day	<i>Apr. 2015</i>
Harvard University, Theory Seminar	<i>Nov. 2014</i>

Service

Organization

NYU Theoretical Computer Science Seminar	<i>2022 - 2025</i>
ICERM at Brown, Working Group Organizer	<i>2026</i>
Simons Institute at Berkeley, Workshop Organizer	<i>2025</i>
SIAM Mathematics of Data Science, Minisymposium Organizer	<i>2020, 2024</i>
NYU Computer Science Theory "Pandemic Presentations" Workshop	<i>2022</i>
SIAM Annual Meeting, Minisymposium Organizer	<i>2021</i>

Department Service

Faculty Search Committee	<i>2019, 2021-2025</i>
Chair, Ph.D. Admissions Committee	<i>2021 - 2023</i>
Department Chair Search Committee	<i>2021 - 2023</i>
Online Faculty Oversight Committee	<i>2021 - 2023</i>

Conference Programming

SIAM LA27 Organizing Committee Member	<i>2026-2027</i>
SODA Program Committee Member	<i>2026</i>
FOCS Program Committee Member	<i>2022, 2025</i>
STOC Program Committee Member	<i>2023</i>
ESA Program Committee Member	<i>2024, 2026</i>
SOSA Program Committee Member	<i>2025</i>
NeurIPS Area Chair	<i>2019, 2021-2024</i>
ICLR Area Chair	<i>2024-2025</i>
ICML Area Chair	<i>2023</i>
AAAI Senior Area Chair	<i>2022</i>

Journal Editing

Editor, SICOMP Special Issue for STOC	<i>2023</i>
---------------------------------------	-------------

Conference Reviewing

STOC, FOCS, SODA, ITCS, ICML, NeurIPS, COLT, AAI, AISTATS, VDLB, ICALP, SOCG, ESA, APPROX, RANDOM, STACS, IWOCA, WALCOM, NDSS,

Journal Reviewing.....

Journal of the ACM, SIAM Journal on Computing, Communications on Pure and Applied Mathematics, SIAM Journal on Matrix Analysis and Applications, SIAM Journal on Scientific Computing, ACM Transactions on Algorithms, Theoretical Computer Science, Journal of Machine Learning Research, Linear Algebra and Its Applications, SIAM Journal on Mathematics of Data Science, Science Advances, Mathematical Programming, IEEE Transactions on Information Theory, IEEE Transactions on Signal Processing, IEEE Transactions on Knowledge and Data Engineering, International Journal of Foundations of Computer Science, IEEE Journal on Selected Areas in Information Theory, BIT Numerical Mathematics, AMS Notices, Bernoulli, International Mathematics Research Notices

Grants Panels.....

National Science Foundation Panel	2019, 2022
NSF Graduate Research Fellowship Reviewer	2022
Department of Energy External Reviewer	2022, 2023, 2024

Honors and Awards

NYU Tandon	2023
<i>Junior Faculty Research Award</i>	
NYU Tandon	2023
<i>Goddard Junior Faculty Fellowship Award</i>	
NYU Tandon	2022
<i>Nominated for Jacobs Excellence in Education Innovation Award</i>	
National Science Foundation	2014-2017
<i>Graduate Research Fellowship, recipient</i>	

Other Professional Experience

Redfin	Seattle, WA
<i>Software Engineer, Data Team</i>	2012 – 2014
Elysium Digital	Cambridge, MA
<i>Summer Consultant</i>	2011
Amicus	New Haven, CT
<i>Software Developer</i>	2010 – 2011