

Christopher Musco

370 Jay Street, Floor 11 – Brooklyn, NY 11201 – Office 1105

☎ (646) 997 3346 • ✉ cmusco@nyu.edu • 🌐 www.chrismusco.com

Academic Positions

New York University <i>Assistant Professor, Tandon School of Engineering</i> <i>Department of Computer Science and Engineering</i>	New York, NY 2019 - present
Princeton University <i>Research Instructor, Department of Computer Science</i>	Princeton, NJ 2018 – 2019
Massachusetts Institute of Technology <i>Research and Teaching Assistant, Theory of Computation Group</i>	Cambridge, MA 2013 – 2018
Yahoo Labs <i>Research Intern, Scalable Machine Learning Group</i>	New York, NY Summer 2015

Research Areas

foundations of data science • scalable machine learning • numerical linear algebra • theory of algorithms
randomization • dimensionality reduction • sketching & streaming • spectral graph algorithms • optimization

Education

Massachusetts Institute of Technology <i>Ph.D. in Computer Science</i> Advisor: Jonathan Kelner Thesis: Faster Linear Algebra for Data Analysis and Machine Learning M.S. Thesis: Dimensionality Reduction for Sparse and Structured Matrices	Cambridge, MA 2013 – 2018
Yale University <i>B.S. in Applied Mathematics, B.S. in Computer Science</i>	New Haven, CT 2008 – 2012

Publications

*Authors appear in alphabetical order, in the tradition of mathematics and theoretical computer science. Any exceptions are marked with a *, which indicates that authors are ordered by contribution.*

Towards Equitable Transit Planning*. Sheng Wang, Yuan Sun, Christopher Musco and Zhifeng Bao. *In submission*. 2020.

Finding the Mode of a Kernel Density Estimate. Jasper C.H. Lee, Jerry Li, Christopher Musco, Jeff M. Phillips, Wai Ming Tai. *In submission*. 2019.

Low-Rank Toeplitz Matrix Estimation via Random Ultra-Sparse Rulers. Hannah Lawrence, Jerry Li, Cameron Musco, Christopher Musco. *International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*. 2020.

Sample Efficient Toeplitz Covariance Estimation. Yonina Eldar, Jerry Li, Cameron Musco, Christopher Musco. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2020.

Fast and Space Efficient Spectral Sparsification in Dynamic Streams. Michael Kapralov, Aida Mousavifar, Cameron Musco, Christopher Musco, Navid Nouri, Aaron Sidford, Aaron Sidford. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2020.

Understanding Filter Bubbles and Polarization in Social Networks. Uthsav Chitra, Christopher Musco. *ACM International Conference on Web Search and Data Mining (WSDM)* 2020.

Low-Rank Approximation from Communication Complexity. Cameron Musco, Christopher Musco, David Woodruff. *In submission.* 2019.

A Universal Sampling Method for Reconstructing Signals with Simple Fourier Transforms. Haim Avron, Michael Kapralov, Cameron Musco, Christopher Musco, Ameya Velingker, Amir Zandieh. *ACM Symposium on Theory of Computing (STOC)* 2019.

Learning Networks from Random Walk-Based Node Similarities. Jeremy Hoskins, Cameron Musco, Christopher Musco, Charalampos Tsourakakis. *Conference on Neural Information Processing Systems (NeurIPS)* 2018.

Eigenvector Computation and Community Detection in Asynchronous Gossip Models. Frederik Mallmann-Trenn, Cameron Musco, Christopher Musco. *International Colloquium on Automata, Languages, and Programming (ICALP)* 2018.

Minimizing Polarization and Disagreement in Social Networks. Cameron Musco, Christopher Musco, Charalampos Tsourakakis. *The Web Conference (WWW)* 2018.

Stability of the Lanczos Method for Matrix Function Approximation. Cameron Musco, Christopher Musco, Aaron Sidford. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2018.

Recursive Sampling for the Nyström Method. Cameron Musco, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2017.

Random Fourier Features for Kernel Ridge Regression: Approximation Bounds and Statistical Guarantees. Haim Avron, Michael Kapralov, Cameron Musco, Christopher Musco, Ameya Velingker, Amir Zandieh. *International Conference on Machine Learning (ICML)* 2017.

Input Sparsity Time Low-Rank Approximation via Ridge Leverage Score Sampling. Michael B. Cohen, Cameron Musco, Christopher Musco. *ACM-SIAM Symposium on Discrete Algorithms (SODA)* 2017.

Determining Tournament Payout Structures for Daily Fantasy Sports. Christopher Musco, Maxim Sviridenko, Justin Thaler. *SIAM Algorithm Engineering and Experiments (ALENEX)* 2017.

Principal Component Projection Without Principal Component Analysis. Roy Frostig, Cameron Musco, Christopher Musco, Aaron Sidford. *International Conference on Machine Learning (ICML)* 2016.

Randomized Block Krylov Methods for Stronger and Faster Approximate Singular Value Decomposition. Cameron Musco, Christopher Musco. *Conference on Neural Information Processing Systems (NeurIPS)* 2015. Invited for full oral presentation (1 of 15 out of 403 accepted papers). Also presented at the 2016 Copper Mountain Conference on Iterative Methods

Dimensionality Reduction for K-Means Clustering and Low Rank Approximation. Michael B. Cohen, Samuel Elder, Cameron Musco, Christopher Musco, Madalina Persu. *ACM Symposium on Theory of Computing (STOC)* 2015.

Principled Sampling for Anomaly Detection. Brendan Juba, Fan Long, Christopher Musco, Stelios Sidiroglou-Douskos, Martin Rinard. *Network and Distributed System Security Symposium (NDSS)* 2015.

Uniform Sampling for Matrix Approximation. Michael B. Cohen, Yin Tat Lee, Cameron Musco, Christopher Musco, Richard Peng, Aaron Sidford. *Innovations in Theoretical Computer Science (ITCS)* 2015.

Single Pass Spectral Sparsification in Dynamic Streams. Michael Kapralov, Yin Tat Lee, Cameron Musco, Christopher Musco, Aaron Sidford. *IEEE Symposium on Foundations of Computer Science (FOCS)* 2014, Special Issue SIAM Journal on Computing 2017.

Student Research Advising

Raphael Meyer. Ph.D., New York University.	2019 - present
Prathamesh Dharangutte. M.S., New York University.	2019 - present
John Shin. M.S., New York University.	2019 - present
Tristan Pollner. Undergraduate, Princeton.	2018 - 2019

Teaching

Introduction to Machine Learning (NYU CS-UY 4563) <i>Course Instructor</i>	Spring 2020
Algorithmic Machine Learning and Data Science (NYU CS-GY 9223I) <i>Course Instructor</i>	Fall 2019
Advanced Algorithm Design (Princeton COS 521) <i>Course Instructor</i>	Fall 2018
Advanced Algorithms (MIT 6.854/18.415) <i>Teaching Assistant</i>	Spring 2016
Technical Communication Skills for Graduate Students (MIT 6.S977) <i>Workshop Leader</i>	Spring 2016
Introduction to Computer Science (Yale CPSC-201) <i>Teaching Assistant</i>	Spring 2012
Introduction to Programming (Yale CPSC-112) <i>Teaching Assistant</i>	Fall 2011

Mentorship and Outreach

SIAM/MathWorks Math Modeling Challenge <i>Problem Development Committee Member, Judge</i>	2018 – present
Support SIAM's annual math modeling competition, which gives high school students an opportunity to apply mathematics, statistics, and data science to a real-world problem. Serve on committee in charge of writing practice and contest problems. Score and provide feedback on student submissions to the contest.	
MIT Graduate Communication Lab <i>Communication Adviser</i>	2015 – 2018
Helped to establish an EECS Communication Lab to provide over 700 graduate students with support in academic writing, oral presentation, job applications, and generally, in communicating effectively. Offered feedback, guidance, and coaching to individual students during weekly office hours.	
CSAIL Algorithms Office Hours <i>Member</i>	2016 – 2018
Advised researchers, generally in applied fields, on framing and solving algorithmic problems arising in their work.	
Research Science Institute at MIT <i>Summer Research Adviser</i>	2014
Primary supervisor for a high school student participating in the prestigious RSI summer research program. Advisee's project, "Relaxation of a Concurrent Disjoint-Set", was selected as one of 10 program finalists.	

Talks and Presentations

Structured Covariance Estimation Johns Hopkins University, Theory Seminar	Oct. 2019
---	-----------

Symposium on Discrete Algorithms (SODA)	<i>Jan. 2020</i>
From Randomized Linear Algebra to Randomized Functional Analysis	
Rutgers University, DIMACS RandNLA Workshop	<i>Sept. 2019</i>
University of Utah, Data Science Seminar	<i>Jan. 2020</i>
Learning Signals with Simple Fourier transforms	
Symposium on Theory of Computing (STOC)	<i>June 2019</i>
Columbia University, Theory Seminar	<i>Oct. 2018</i>
UC Berkeley, Simons Institute Workshop	<i>Sept. 2018</i>
Princeton University, Theory Seminar	<i>Sept. 2018</i>
The Lanczos Method in Data Science: Challenges and Opportunities	
SIAM Conference on Parallel Processing for Scientific Computing, Minisymposium	<i>Feb. 2020</i>
Georgetown University, CS Seminar	<i>Feb. 2018</i>
Symposium on Discrete Algorithms (SODA)	<i>Jan. 2018</i>
Unifying Random Features and Leverage Scores for Kernel Matrix Approximation	
Institute for Advanced Study, Computer Science and Discrete Mathematics Seminar	<i>Feb. 2018</i>
Harvard University, guest lecture in "Sketching Algorithms for Big Data"	<i>Nov. 2017</i>
Microsoft Research, Redmond	<i>Sept. 2017</i>
Sketching and Sampling Methods for Large Scale Linear Algebra	
Twitter, Cambridge	<i>July 2017</i>
Recursive Sampling for the Nyström Method	
Neural Information Processing Systems (NeurIPS) poster session	<i>Dec. 2017</i>
New England Machine Learning Day poster session	<i>May 2017</i>
Algorithms for Determining Tournament Payout Structures	
Meeting on Algorithm Engineering and Experiments (ALENEX)	<i>Jan. 2017</i>
Introduction to Linear Sketching	
MIT, Theory Retreat	<i>Sept. 2016</i>
Principal Component Projection without Principal Component Analysis	
International Conference on Machine Learning (ICML)	<i>June 2016</i>
Ridge Leverage Score Sampling	
National Institute of Informatics, Shonan Meeting	<i>July 2016</i>
MIT, Algorithms & Complexity Seminar	<i>July 2016</i>
University of Utah, Data Group Meeting	<i>Jan. 2016</i>
Randomized Block Krylov Methods	
Neural Information Processing Systems (NeurIPS) poster session	<i>Dec. 2015</i>
Large Scale Column Subset Selection	
Yahoo Labs, Science Week poster session	<i>July 2015</i>
Dimensionality Reduction for k-Means Clustering	
IBM T.J. Watson, Mathematical Sciences research seminar	<i>Aug. 2015</i>
Symposium on Theory of Computing (STOC)	<i>June 2015</i>
MIT, Theory Lunch	<i>June 2014</i>
Principled Sampling for Anomaly Detection	

MIT, Defense Advanced Research Projects Agency site visit	Mar. 2015
Network and Distributed System Security Symposium (NDSS)	Feb. 2015
Uniform Sampling for Matrix Approximation	
Yahoo Labs, Scalable Machine Learning research seminar	June 2015
Single Pass Spectral Sparsification in Dynamic Streams	
MIT, Annual Sublinear Algorithms Day poster session	Apr. 2015
Harvard University, Theory Seminar	Nov. 2014
Foundations of Computer Science (FOCS)	Oct. 2014

Service

Conference Programming.....	
NeuriPS Area Chair	2019
Organization.....	
SIAM Mathematics of Data Science Minisymposium Organizer	2020
NYU Tandon Algorithmic Data Science Reading Group	2019 - present
Department Service.....	
Faculty Search Committee	2019
Ph.D. Admissions Committee	2019 - present
Conference Reviewing.....	
STOC, FOCS, SODA, ITCS, ICALP, ESA, ICML, NeuRIPS, APPROX, RANDOM, STACS, IWOC, WALCOM, NDSS	
Journal Reviewing.....	
Journal of Machine Learning Research, SIAM Journal on Matrix Analysis and Applications, SIAM Journal on Scientific Computing, SIAM Journal on Mathematics of Data Science, Science Advances, Mathematical Programming, IEEE Transactions on Signal Processing, Advances in Data Analysis and Classification, IEEE Transactions on Knowledge and Data Engineering, International Journal of Foundations of Computer Science.	
Grants Panels.....	
NSF Division of Information and Intelligent Systems (2019)	
Notables.....	
ICML 2018 Outstanding Reviewer	

Other Research Experience

Rutgers, DIMACS	
<i>Invited participant</i>	2019
Workshop on Randomized Numerical Linear Algebra, Statistics, and Optimization.	
UC Berkeley, Simons Institute for the Theory of Computing	
<i>Invited participant</i>	2018
Workshop on Randomized Numerical Linear Algebra and Applications.	
Japan National Institute of Informatics	
<i>Invited participant</i>	2016
Shonan meeting on Recent Advances in Randomized Numerical Linear Algebra.	

Honors and Awards

National Science Foundation

2014-2017

Graduate Research Fellowship, recipient.

Yale University

2012

Cum Laude, Distinction in Both Majors.

Other Professional Experience

Redfin

Seattle, WA

Software Engineer, Data Team

2012 – 2014

Developed backend Java infrastructure for internet-powered real estate startup. Incorporated large-scale machine learning algorithms into production code. Managed three interns and onboarded four new-hires.

Elysium Digital

Cambridge, MA

Summer Consultant

2011

Provided technical expertise in software patent cases at leading litigation consulting company.

Amicus

New Haven, CT

Software Developer

2010 – 2011

Built applications for political campaign management and social fundraising at an early-stage startup that went on to raise 4 million dollars in seed funding.